

Lending Club Loan Status Prediction - Neural Network Classification

```
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import classification_report, confusion_matrix,
roc_auc_score, roc_curve
from sklearn.impute import SimpleImputer
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import warnings
warnings.filterwarnings('ignore')

# Set random seed for reproducibility
np.random.seed(42)
tf.random.set_seed(42)

# Configure matplotlib and seaborn
plt.style.use('default')
sns.set_palette("husl")

print("All libraries imported successfully!")
print(f"TensorFlow version: {tf.__version__}")

All libraries imported successfully!
TensorFlow version: 2.18.0
```

1. DATA LOADING AND INITIAL EXPLORATION

```
# Load the dataset
try:
    df = pd.read_csv('lending_club_loan_two.csv')
    print("Dataset loaded successfully!")
    print(f"Dataset shape: {df.shape}")
except FileNotFoundError:
    print("Please upload your 'lending_club_loan_two.csv' file to Colab first")
    print("Use the file upload button in the sidebar or:")
```

```
print("from google.colab import files")
print("files.upload()")
```

Dataset loaded successfully!
Dataset shape: (396030, 27)

Display basic information about the dataset

```
print("DATASET OVERVIEW")
print(f"Dataset shape: {df.shape}")
print(f"Memory usage: {df.memory_usage(deep=True).sum() / 1024**2:.2f} MB")
```

DATASET OVERVIEW

Dataset shape: (396030, 27)

Memory usage: 421.42 MB

Display first few rows

```
print("\nFirst 5 rows:")
print(df.head())
```

First 5 rows:

	loan_amnt	term	int_rate	installment	grade	sub_grade	\
0	10000.0	36 months	11.44	329.48	B	B4	
1	8000.0	36 months	11.99	265.68	B	B5	
2	15600.0	36 months	10.49	506.97	B	B3	
3	7200.0	36 months	6.49	220.65	A	A2	
4	24375.0	60 months	17.27	609.33	C	C5	

	emp_title	emp_length	home_ownership	annual_inc	...
0	Marketing	10+ years	RENT	117000.0	...
1	Credit analyst	4 years	MORTGAGE	65000.0	...
2	Statistician	< 1 year	RENT	43057.0	...
3	Client Advocate	6 years	RENT	54000.0	...
4	Destiny Management Inc.	9 years	MORTGAGE	55000.0	...

	open_acc	pub_rec	revol_bal	revol_util	total_acc	initial_list_status
0	16.0	0.0	36369.0	41.8	25.0	w
1	17.0	0.0	20131.0	53.3	27.0	f
2	13.0	0.0	11987.0	92.2	26.0	f
3	6.0	0.0	5472.0	21.5	13.0	f

4	13.0	0.0	24584.0	69.8	43.0	f
---	------	-----	---------	------	------	---

	application_type	mort_acc	pub_rec_bankruptcies	\
0	INDIVIDUAL	0.0	0.0	
1	INDIVIDUAL	3.0	0.0	
2	INDIVIDUAL	0.0	0.0	
3	INDIVIDUAL	0.0	0.0	
4	INDIVIDUAL	1.0	0.0	

	address
0	0174 Michelle Gateway\r\nMendozaberg, OK 22690
1	1076 Carney Fort Apt. 347\r\nLoganmouth, SD 05113
2	87025 Mark Dale Apt. 269\r\nNew Sabrina, WV 05113
3	823 Reid Ford\r\nDelacruzside, MA 00813
4	679 Luna Roads\r\nGreggshire, VA 11650

[5 rows x 27 columns]

```
# Display data types and missing values
print("\nData Types and Missing Values:")
info_df = pd.DataFrame({
    'Column': df.columns,
    'Data Type': df.dtypes,
    'Missing Values': df.isnull().sum(),
    'Missing %': (df.isnull().sum() / len(df) * 100).round(2)
})
print(info_df)
```

Data Types and Missing Values:

	Column	Data Type	Missing
Values \			
loan_amnt	loan_amnt	float64	0
term	term	object	0
int_rate	int_rate	float64	0
installment	installment	float64	0
grade	grade	object	0
sub_grade	sub_grade	object	0
emp_title	emp_title	object	22927
emp_length	emp_length	object	18301
home_ownership	home_ownership	object	0

annual_inc	annual_inc	float64	0
verification_status	verification_status	object	0
issue_d	issue_d	object	0
loan_status	loan_status	object	0
purpose	purpose	object	0
title	title	object	1756
dti	dti	float64	0
earliest_cr_line	earliest_cr_line	object	0
open_acc	open_acc	float64	0
pub_rec	pub_rec	float64	0
revol_bal	revol_bal	float64	0
revol_util	revol_util	float64	276
total_acc	total_acc	float64	0
initial_list_status	initial_list_status	object	0
application_type	application_type	object	0
mort_acc	mort_acc	float64	37795
pub_rec_bankruptcies	pub_rec_bankruptcies	float64	535
address	address	object	0

	Missing %
loan_amnt	0.00
term	0.00
int_rate	0.00
installment	0.00
grade	0.00
sub_grade	0.00
emp_title	5.79
emp_length	4.62
home_ownership	0.00
annual_inc	0.00
verification_status	0.00
issue_d	0.00

loan_status	0.00
purpose	0.00
title	0.44
dti	0.00
earliest_cr_line	0.00
open_acc	0.00
pub_rec	0.00
revol_bal	0.00
revol_util	0.07
total_acc	0.00
initial_list_status	0.00
application_type	0.00
mort_acc	9.54
pub_rec_bankruptcies	0.14
address	0.00

Basic statistics

print("\nBasic Statistics for Numerical Columns:")

print(df.describe())

Basic Statistics for Numerical Columns:

	loan_amnt	int_rate	installment	annual_inc	\
count	396030.000000	396030.000000	396030.000000	3.960300e+05	
mean	14113.888089	13.639400	431.849698	7.420318e+04	
std	8357.441341	4.472157	250.727790	6.163762e+04	
min	500.000000	5.320000	16.080000	0.000000e+00	
25%	8000.000000	10.490000	250.330000	4.500000e+04	
50%	12000.000000	13.330000	375.430000	6.400000e+04	
75%	20000.000000	16.490000	567.300000	9.000000e+04	
max	40000.000000	30.990000	1533.810000	8.706582e+06	

	dti	open_acc	pub_rec	revol_bal	\
count	396030.000000	396030.000000	396030.000000	3.960300e+05	
mean	17.379514	11.311153	0.178191	1.584454e+04	
std	18.019092	5.137649	0.530671	2.059184e+04	
min	0.000000	0.000000	0.000000	0.000000e+00	
25%	11.280000	8.000000	0.000000	6.025000e+03	
50%	16.910000	10.000000	0.000000	1.118100e+04	
75%	22.980000	14.000000	0.000000	1.962000e+04	
max	9999.000000	90.000000	86.000000	1.743266e+06	


	revol_util	total_acc	mort_acc
pub_rec_bankruptcies			
count	395754.000000	396030.000000	358235.000000
mean	53.791749	25.414744	1.813991
std	24.452193	11.886991	2.147930
min	0.000000	0.000000	0.000000
25%	11.280000	8.000000	0.000000
50%	16.910000	10.000000	0.000000
75%	22.980000	14.000000	0.000000
max	9999.000000	90.000000	86.000000

min	0.000000	2.000000	0.000000
0.000000			
25%	35.800000	17.000000	0.000000
0.000000			
50%	54.800000	24.000000	1.000000
0.000000			
75%	72.900000	32.000000	3.000000
0.000000			
max	892.300000	151.000000	34.000000
8.000000			

2. EXPLORATORY DATA ANALYSIS (EDA)

```
# Create figure for EDA plots
fig, axes = plt.subplots(2, 2, figsize=(15, 12))
fig.suptitle('Lending Club Dataset - Exploratory Data Analysis',
             fontsize=16)
Text(0.5, 0.98, 'Lending Club Dataset - Exploratory Data Analysis')
```

Lending Club Dataset - Exploratory Data Analysis



```
# 1. Target Variable Distribution
if 'loan_status' in df.columns:
    loan_status_counts = df['loan_status'].value_counts()
    axes[0, 0].pie(loan_status_counts.values,
labels=loan_status_counts.index, autopct='%1.1f%%')
    axes[0, 0].set_title('Distribution of Loan Status')
    print(f"\nLoan Status Distribution:")
    print(loan_status_counts)
    print(f"Class balance ratio:
{loan_status_counts.min()/loan_status_counts.max():.3f}")
```

```
Loan Status Distribution:
loan_status
Fully Paid      318357
Charged Off     77673
```

```
Name: count, dtype: int64
Class balance ratio: 0.244
```

2. Loan Amount Distribution

```
if 'loan_amnt' in df.columns:
    axes[0, 1].hist(df['loan_amnt'].dropna(), bins=50, alpha=0.7,
color='skyblue')
    axes[0, 1].set_title('Distribution of Loan Amount')
    axes[0, 1].set_xlabel('Loan Amount ($)')
    axes[0, 1].set_ylabel('Frequency')
```

3. Interest Rate Distribution

```
if 'int_rate' in df.columns:
    axes[1, 0].hist(df['int_rate'].dropna(), bins=30, alpha=0.7,
color='lightcoral')
    axes[1, 0].set_title('Distribution of Interest Rate')
    axes[1, 0].set_xlabel('Interest Rate (%)')
    axes[1, 0].set_ylabel('Frequency')
```

4. Grade Distribution

```
if 'grade' in df.columns:
    grade_counts = df['grade'].value_counts().sort_index()
    axes[1, 1].bar(grade_counts.index, grade_counts.values,
color='lightgreen')
    axes[1, 1].set_title('Distribution of Loan Grades')
    axes[1, 1].set_xlabel('Grade')
    axes[1, 1].set_ylabel('Count')
```


```
plt.tight_layout()
plt.show()
```

<Figure size 640x480 with 0 Axes>

Correlation Analysis for Numerical Features

Select numerical columns for correlation analysis

```
numerical_cols =
df.select_dtypes(include=[np.number]).columns.tolist()
if len(numerical_cols) > 1:
    plt.figure(figsize=(12, 8))
    correlation_matrix = df[numerical_cols].corr()
    sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm',
center=0,
                square=True, fmt='.2f', cbar_kws={"shrink": .8})
    plt.title('Correlation Matrix of Numerical Features')
    plt.tight_layout()
    plt.show()
```

```
# Find highly correlated features
high_corr_pairs = []
for i in range(len(correlation_matrix.columns)):
    for j in range(i+1, len(correlation_matrix.columns)):
        if abs(correlation_matrix.iloc[i, j]) > 0.7:
            high_corr_pairs.append((
                correlation_matrix.columns[i],
                correlation_matrix.columns[j],
                correlation_matrix.iloc[i, j]
            ))

if high_corr_pairs:
    print("Highly correlated feature pairs (|r| > 0.7):")
    for pair in high_corr_pairs:
        print(f"  {pair[0]} - {pair[1]}: {pair[2]:.3f}")
```

```

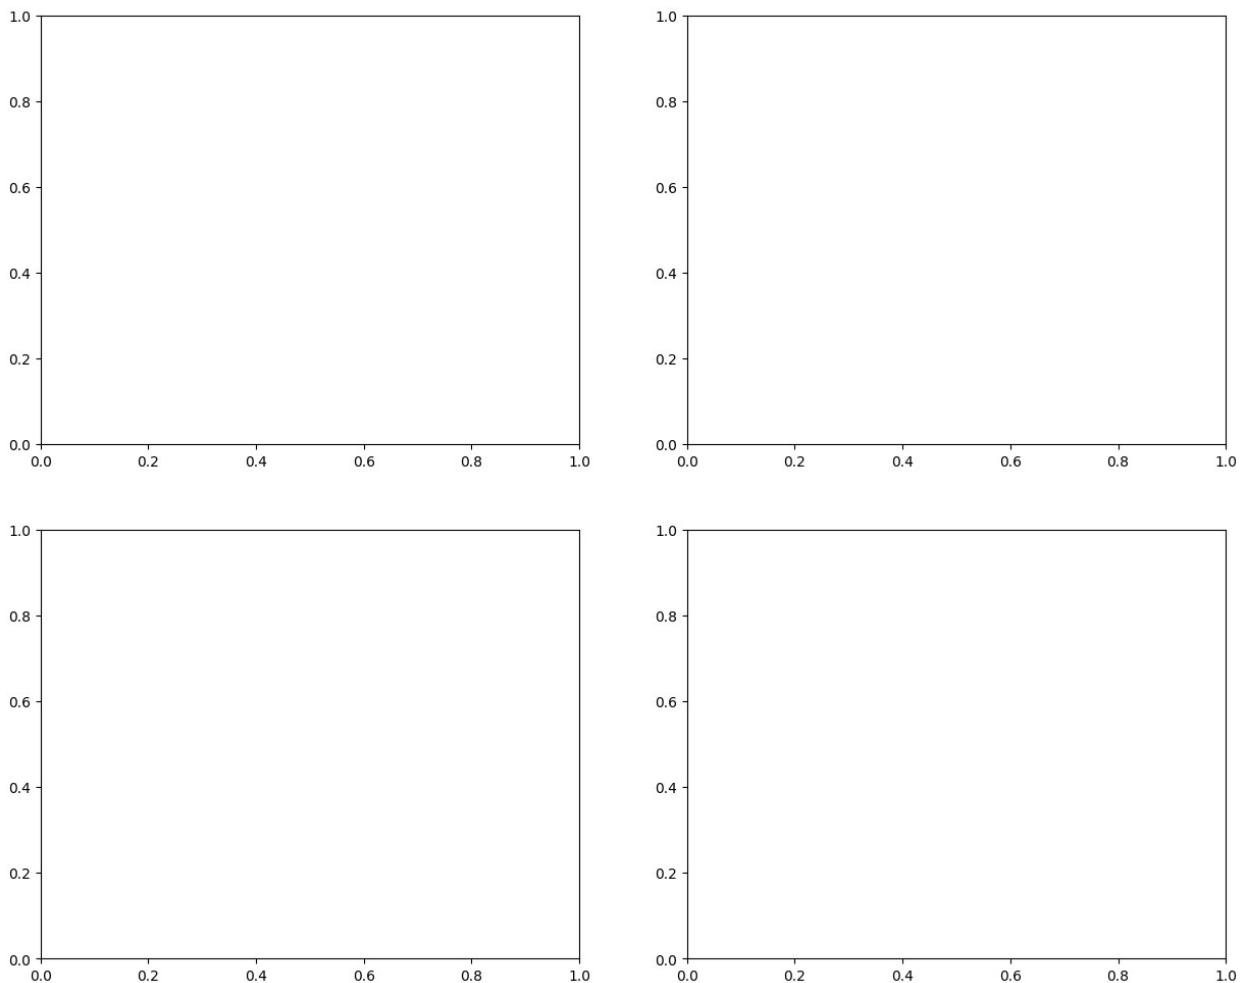
else:
    print("No highly correlated feature pairs found ( $|r| > 0.7$ )")

Highly correlated feature pairs ( $|r| > 0.7$ ):
    loan_amnt - installment: 0.954

# Additional EDA plots
fig, axes = plt.subplots(2, 2, figsize=(15, 12))
fig.suptitle('Additional EDA - Feature Relationships', fontsize=16)
Text(0.5, 0.98, 'Additional EDA - Feature Relationships')

```

Additional EDA - Feature Relationships



```

# Loan Amount vs Interest Rate
if 'loan_amnt' in df.columns and 'int_rate' in df.columns:
    axes[0, 0].scatter(df['loan_amnt'], df['int_rate'], alpha=0.5)
    axes[0, 0].set_xlabel('Loan Amount ($)')

```

```

    axes[0, 0].set_ylabel('Interest Rate (%)')
    axes[0, 0].set_title('Loan Amount vs Interest Rate')

# Employment Length Distribution
if 'emp_length' in df.columns:
    emp_length_counts = df['emp_length'].value_counts()
    axes[0, 1].bar(range(len(emp_length_counts)),
emp_length_counts.values)
    axes[0, 1].set_title('Employment Length Distribution')
    axes[0, 1].set_xlabel('Employment Length')
    axes[0, 1].set_ylabel('Count')
    axes[0, 1].set_xticks(range(len(emp_length_counts)))
    axes[0, 1].set_xticklabels(emp_length_counts.index, rotation=45)

# Home Ownership Distribution
if 'home_ownership' in df.columns:
    home_counts = df['home_ownership'].value_counts()
    axes[1, 0].pie(home_counts.values, labels=home_counts.index,
autopct='%1.1f%%')
    axes[1, 0].set_title('Home Ownership Distribution')

# Annual Income Distribution (log scale)
if 'annual_inc' in df.columns:
    axes[1, 1].hist(np.log1p(df['annual_inc'].dropna()), bins=50,
alpha=0.7, color='gold')
    axes[1, 1].set_title('Distribution of Annual Income (Log Scale)')
    axes[1, 1].set_xlabel('Log(Annual Income + 1)')
    axes[1, 1].set_ylabel('Frequency')

plt.tight_layout()
plt.show()

<Figure size 640x480 with 0 Axes>

```

3. DATA PREPROCESSING

```

# Create a copy of the dataframe for preprocessing
df_processed = df.copy()

# Handle missing values
missing_summary = df_processed.isnull().sum()
print(f"Columns with missing values: {missing_summary[missing_summary
> 0].shape[0]}")

Columns with missing values: 6

# For numerical columns, use median imputation
numerical_cols =
df_processed.select_dtypes(include=[np.number]).columns

```

```

for col in numerical_cols:
    if df_processed[col].isnull().sum() > 0:
        median_val = df_processed[col].median()
        df_processed[col].fillna(median_val, inplace=True)
        print(f"Filled {col} missing values with median:
{median_val}")

Filled revol_util missing values with median: 54.8
Filled mort_acc missing values with median: 1.0
Filled pub_rec_bankruptcies missing values with median: 0.0

# For categorical columns, use mode imputation
categorical_cols =
df_processed.select_dtypes(include=['object']).columns
for col in categorical_cols:
    if df_processed[col].isnull().sum() > 0:
        mode_val = df_processed[col].mode()[0] if not
df_processed[col].mode().empty else 'Unknown'
        df_processed[col].fillna(mode_val, inplace=True)
        print(f"Filled {col} missing values with mode: {mode_val}")

Filled emp_title missing values with mode: Teacher
Filled emp_length missing values with mode: 10+ years
Filled title missing values with mode: Debt consolidation

# Prepare target variable
if 'loan_status' in df_processed.columns:
    # Create binary target variable
    # Assuming 'Fully Paid' = 1 (good), others = 0 (bad)
    df_processed['target'] = (df_processed['loan_status'] == 'Fully
Paid').astype(int)
    print(f"Target variable created. Distribution:")
    print(df_processed['target'].value_counts())

Target variable created. Distribution:
target
1    318357
0     77673
Name: count, dtype: int64

```

Feature Engineering

```

# Extract year from issue_d if it exists
if 'issue_d' in df_processed.columns:
    df_processed['issue_year'] =
pd.to_datetime(df_processed['issue_d'], errors='coerce').dt.year
    print("Extracted issue year from issue_d")

Extracted issue year from issue_d

```

```
# Extract year from earliest_cr_line if it exists
```

```
if 'earliest_cr_line' in df_processed.columns:  
    df_processed['earliest_cr_year'] =  
pd.to_datetime(df_processed['earliest_cr_line'],  
errors='coerce').dt.year  
    print("Extracted earliest credit line year")
```

Extracted earliest credit line year

```
# Calculate credit history length
```

```
if 'issue_year' in df_processed.columns and 'earliest_cr_year' in  
df_processed.columns:  
    df_processed['credit_history_length'] = df_processed['issue_year']  
- df_processed['earliest_cr_year']  
    print("Calculated credit history length")
```

Calculated credit history length

```
# Create debt-to-income ratio categories
```

```
if 'dti' in df_processed.columns:  
    df_processed['dti_category'] = pd.cut(df_processed['dti'],  
bins=[0, 10, 20, 30,  
float('inf')]),  
labels=['Low', 'Medium',  
'High', 'Very High'])  
    print("Created DTI categories")
```

Created DTI categories

```
# Process categorical variables
```

```
print("\nEncoding categorical variables...")
```

```
categorical_features = []
```

```
label_encoders = {}
```

```
for col in categorical_cols:
```

```
    if col not in ['loan_status', 'target']: # Skip target-related  
columns
```

```
        if df_processed[col].nunique() <= 10: # Use label encoding  
for low cardinality
```

```
            le = LabelEncoder()  
            df_processed[f'{col}_encoded'] =  
le.fit_transform(df_processed[col].astype(str))  
            label_encoders[col] = le  
            categorical_features.append(f'{col}_encoded')  
            print(f"Label encoded {col} ({df_processed[col].nunique()}  
unique values)")
```

```
        else: # Use one-hot encoding for high cardinality (top  
categories only)
```

```
            top_categories =  
df_processed[col].value_counts().head(10).index  
            for category in top_categories:
```

```

        df_processed[f'{col}_{category}'] = (df_processed[col]
== category).astype(int)
        categorical_features.append(f'{col}_{category}')
    print(f"One-hot encoded top 10 categories for {col}")

```

Encoding categorical variables...

```

Label encoded term (2 unique values)
Label encoded grade (7 unique values)
One-hot encoded top 10 categories for sub_grade
One-hot encoded top 10 categories for emp_title
One-hot encoded top 10 categories for emp_length
Label encoded home_ownership (6 unique values)
Label encoded verification_status (3 unique values)
One-hot encoded top 10 categories for issue_d
One-hot encoded top 10 categories for purpose
One-hot encoded top 10 categories for title
One-hot encoded top 10 categories for earliest_cr_line
Label encoded initial_list_status (2 unique values)
Label encoded application_type (3 unique values)
One-hot encoded top 10 categories for address

```

Select features for modeling

```
feature_cols = []
```

Add numerical features

```

numerical_features = ['loan_amnt', 'int_rate', 'installment',
'annual_inc', 'dti',
                        'open_acc', 'pub_rec', 'revol_bal', 'revol_util',
'total_acc',
                        'mort_acc', 'pub_rec_bankruptcies']

```

```

for col in numerical_features:
    if col in df_processed.columns:
        feature_cols.append(col)

```

Add engineered features

```

engineered_features = ['issue_year', 'earliest_cr_year',
'credit_history_length']
for col in engineered_features:
    if col in df_processed.columns:
        feature_cols.append(col)

```

Add categorical features

```
feature_cols.extend(categorical_features)
```

Remove any features that don't exist in the dataset

```
feature_cols = [col for col in feature_cols if col in
df_processed.columns]
```

```
print(f"Selected {len(feature_cols)} features for modeling")
print(f"Features: {feature_cols}")
```

Selected 101 features for modeling

```
Features: ['loan_amnt', 'int_rate', 'installment', 'annual_inc',
'dti', 'open_acc', 'pub_rec', 'revol_bal', 'revol_util', 'total_acc',
'mort_acc', 'pub_rec_bankruptcies', 'issue_year', 'earliest_cr_year',
'credit_history_length', 'term_encoded', 'grade_encoded',
'sub_grade_B3', 'sub_grade_B4', 'sub_grade_C1', 'sub_grade_C2',
'sub_grade_B2', 'sub_grade_B5', 'sub_grade_C3', 'sub_grade_C4',
'sub_grade_B1', 'sub_grade_A5', 'emp_title_Teacher',
'emp_title_Manager', 'emp_title_Registered Nurse', 'emp_title_RN',
'emp_title_Supervisor', 'emp_title_Sales', 'emp_title_Project
Manager', 'emp_title_Owner', 'emp_title_Driver', 'emp_title_Office
Manager', 'emp_length_10+ years', 'emp_length_2 years', 'emp_length_<
1 year', 'emp_length_3 years', 'emp_length_5 years', 'emp_length_1
year', 'emp_length_4 years', 'emp_length_6 years', 'emp_length_7
years', 'emp_length_8 years', 'home_ownership_encoded',
'verification_status_encoded', 'issue_d_Oct-2014', 'issue_d_Jul-2014',
'issue_d_Jan-2015', 'issue_d_Dec-2013', 'issue_d_Nov-2013',
'issue_d_Jul-2015', 'issue_d_Oct-2013', 'issue_d_Jan-2014',
'issue_d_Apr-2015', 'issue_d_Sep-2013', 'purpose_debt_consolidation',
'purpose_credit_card', 'purpose_home_improvement', 'purpose_other',
'purpose_major_purchase', 'purpose_small_business', 'purpose_car',
'purpose_medical', 'purpose_moving', 'purpose_vacation', 'title_Debt
consolidation', 'title_Credit card refinancing', 'title_Home
improvement', 'title_Other', 'title_Debt Consolidation', 'title_Major
purchase', 'title_Consolidation', 'title_debt consolidation',
'title_Business', 'title_Debt Consolidation Loan',
'earliest_cr_line_Oct-2000', 'earliest_cr_line_Aug-2000',
'earliest_cr_line_Oct-2001', 'earliest_cr_line_Aug-2001',
'earliest_cr_line_Nov-2000', 'earliest_cr_line_Oct-1999',
'earliest_cr_line_Nov-1999', 'earliest_cr_line_Sep-2000',
'earliest_cr_line_Oct-2002', 'earliest_cr_line_Aug-2002',
'initial_list_status_encoded', 'application_type_encoded',
'address_USS Johnson\r\nFPO AE 48052', 'address_USNS Johnson\r\nFPO AE
05113', 'address_USS Smith\r\nFPO AP 70466', 'address_USCGC Smith\r\
nFPO AE 70466', 'address_USNS Johnson\r\nFPO AP 48052', 'address_USNV
Smith\r\nFPO AA 00813', 'address_USNV Smith\r\nFPO AE 30723',
'address_USCGC Miller\r\nFPO AA 22690', 'address_USNS Johnson\r\nFPO
AA 70466', 'address_USS Smith\r\nFPO AP 22690']
```

```
# Prepare final dataset
```

```
if 'target' in df_processed.columns:
```

```
    # Remove rows with missing target
```

```
    df_final = df_processed.dropna(subset=['target'])
```

```
    X = df_final[feature_cols]
```

```
    y = df_final['target']
```

```

    print(f"Final dataset shape: {X.shape}")
    print(f"Target distribution: {y.value_counts().to_dict()}")
else:
    print("Warning: Target variable 'loan_status' not found. Creating dummy target.")
    X = df_processed[feature_cols]
    y = np.random.choice([0, 1], size=len(X)) # Dummy target for demonstration

Final dataset shape: (396030, 101)
Target distribution: {1: 318357, 0: 77673}

# Handle any remaining missing values
X = X.fillna(X.median())

```

4. MODEL BUILDING

```

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42, stratify=y)

print(f"Training set shape: {X_train.shape}")
print(f"Testing set shape: {X_test.shape}")
print(f"Training set target distribution:
{pd.Series(y_train).value_counts().to_dict()}")
print(f"Testing set target distribution:
{pd.Series(y_test).value_counts().to_dict()}")

Training set shape: (316824, 101)
Testing set shape: (79206, 101)
Training set target distribution: {1: 254686, 0: 62138}
Testing set target distribution: {1: 63671, 0: 15535}

# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

Features scaled using StandardScaler

# Build Neural Network Model
model = keras.Sequential([
    layers.Dense(128, activation='relu',
input_shape=(X_train_scaled.shape[1],)),
    layers.Dropout(0.3),
    layers.Dense(64, activation='relu'),
    layers.Dropout(0.3),
    layers.Dense(32, activation='relu'),
    layers.Dropout(0.2),

```



```

        layers.Dense(16, activation='relu'),
        layers.Dense(1, activation='sigmoid')
    ])

# Compile the model
model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy', 'precision', 'recall']
)

# Display model architecture
print("\nModel Architecture:")
model.summary()

```

Model Architecture:

Model: "sequential"

Layer (type) Param #	Output Shape
dense (Dense) 13,056	(None, 128)
dropout (Dropout) 0	(None, 128)
dense_1 (Dense) 8,256	(None, 64)
dropout_1 (Dropout) 0	(None, 64)
dense_2 (Dense) 2,080	(None, 32)
dropout_2 (Dropout) 0	(None, 32)

528	dense_3 (Dense)	(None, 16)	
17	dense_4 (Dense)	(None, 1)	

Total params: 23,937 (93.50 KB)

Trainable params: 23,937 (93.50 KB)

Non-trainable params: 0 (0.00 B)

```
# Train the model
print("\nTraining the model...")
history = model.fit(
    X_train_scaled, y_train,
    epochs=30,
    batch_size=32,
    validation_split=0.2,
    verbose=1,
    callbacks=[
        keras.callbacks.EarlyStopping(monitor='val_loss', patience=10,
        restore_best_weights=True),
        keras.callbacks.ReduceLROnPlateau(monitor='val_loss',
        factor=0.5, patience=5, min_lr=0.0001)
    ]
)
```

Training the model...

Epoch 1/30

7921/7921 ————— 30s 4ms/step - accuracy: 0.8124 - loss: 0.4339 - precision: 0.8219 - recall: 0.9790 - val_accuracy: 0.8063 - val_loss: 0.4458 - val_precision: 0.8117 - val_recall: 0.9883 - learning_rate: 1.2500e-04

Epoch 2/30

7921/7921 ————— 30s 4ms/step - accuracy: 0.8130 - loss: 0.4333 - precision: 0.8228 - recall: 0.9782 - val_accuracy: 0.8064 - val_loss: 0.4456 - val_precision: 0.8119 - val_recall: 0.9881 - learning_rate: 1.2500e-04

Epoch 3/30

7921/7921 ————— 39s 4ms/step - accuracy: 0.8135 - loss: 0.4330 - precision: 0.8235 - recall: 0.9778 - val_accuracy: 0.8065 - val_loss: 0.4458 - val_precision: 0.8120 - val_recall: 0.9880 - learning_rate: 1.2500e-04

Epoch 4/30

7921/7921 ————— 31s 4ms/step - accuracy: 0.8131 - loss:

0.4328 - precision: 0.8231 - recall: 0.9779 - val_accuracy: 0.8065 -
val_loss: 0.4459 - val_precision: 0.8121 - val_recall: 0.9878 -
learning_rate: 1.2500e-04
Epoch 5/30
7921/7921 ————— 38s 3ms/step - accuracy: 0.8125 - loss:
0.4337 - precision: 0.8224 - recall: 0.9782 - val_accuracy: 0.8064 -
val_loss: 0.4459 - val_precision: 0.8115 - val_recall: 0.9889 -
learning_rate: 1.2500e-04
Epoch 6/30
7921/7921 ————— 29s 4ms/step - accuracy: 0.8125 - loss:
0.4333 - precision: 0.8225 - recall: 0.9780 - val_accuracy: 0.8063 -
val_loss: 0.4459 - val_precision: 0.8116 - val_recall: 0.9885 -
learning_rate: 1.2500e-04
Epoch 7/30
7921/7921 ————— 28s 4ms/step - accuracy: 0.8127 - loss:
0.4333 - precision: 0.8225 - recall: 0.9783 - val_accuracy: 0.8066 -
val_loss: 0.4459 - val_precision: 0.8119 - val_recall: 0.9885 -
learning_rate: 1.2500e-04
Epoch 8/30
7921/7921 ————— 43s 4ms/step - accuracy: 0.8137 - loss:
0.4331 - precision: 0.8236 - recall: 0.9777 - val_accuracy: 0.8063 -
val_loss: 0.4458 - val_precision: 0.8120 - val_recall: 0.9877 -
learning_rate: 1.0000e-04
Epoch 9/30
7921/7921 ————— 39s 4ms/step - accuracy: 0.8130 - loss:
0.4326 - precision: 0.8229 - recall: 0.9781 - val_accuracy: 0.8063 -
val_loss: 0.4460 - val_precision: 0.8127 - val_recall: 0.9864 -
learning_rate: 1.0000e-04
Epoch 10/30
7921/7921 ————— 30s 4ms/step - accuracy: 0.8130 - loss:
0.4332 - precision: 0.8231 - recall: 0.9778 - val_accuracy: 0.8064 -
val_loss: 0.4456 - val_precision: 0.8125 - val_recall: 0.9870 -
learning_rate: 1.0000e-04
Epoch 11/30
7921/7921 ————— 29s 4ms/step - accuracy: 0.8131 - loss:
0.4325 - precision: 0.8232 - recall: 0.9777 - val_accuracy: 0.8066 -
val_loss: 0.4459 - val_precision: 0.8127 - val_recall: 0.9868 -
learning_rate: 1.0000e-04
Epoch 12/30
7921/7921 ————— 28s 4ms/step - accuracy: 0.8130 - loss:
0.4325 - precision: 0.8234 - recall: 0.9772 - val_accuracy: 0.8063 -
val_loss: 0.4460 - val_precision: 0.8121 - val_recall: 0.9874 -
learning_rate: 1.0000e-04

5. MODEL EVALUATION

```
# Make predictions
y_pred_prob = model.predict(X_test_scaled)
y_pred = (y_pred_prob > 0.5).astype(int)

2476/2476 ————— 4s 2ms/step

# Calculate metrics
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score

accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc_score = roc_auc_score(y_test, y_pred_prob)

print("Model Performance Metrics:")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")
print(f"AUC-ROC: {auc_score:.4f}")

Model Performance Metrics:
Accuracy: 0.8066
Precision: 0.8117
Recall: 0.9887
F1-Score: 0.8915
AUC-ROC: 0.7250

# Detailed classification report
print("\nDetailed Classification Report:")
print(classification_report(y_test, y_pred))

Detailed Classification Report:

```

	precision	recall	f1-score	support
0	0.57	0.06	0.11	15535
1	0.81	0.99	0.89	63671
accuracy			0.81	79206
macro avg	0.69	0.52	0.50	79206
weighted avg	0.76	0.81	0.74	79206

```
# Confusion Matrix
print("\nConfusion Matrix:")
```

```
cm = confusion_matrix(y_test, y_pred)
print(cm)
```

Confusion Matrix:

```
[[ 934 14601]
 [ 718 62953]]
```


```
# Visualization of results
```

```
fig, axes = plt.subplots(2, 2, figsize=(15, 12))
```

```
fig.suptitle('Model Evaluation Results', fontsize=16)
```

```
Text(0.5, 0.98, 'Model Evaluation Results')
```

Model Evaluation Results



```
# Training history
```

```
axes[0, 0].plot(history.history['accuracy'], label='Training')
```

```

Accuracy')
axes[0, 0].plot(history.history['val_accuracy'], label='Validation
Accuracy')
axes[0, 0].set_title('Model Accuracy')
axes[0, 0].set_xlabel('Epoch')
axes[0, 0].set_ylabel('Accuracy')
axes[0, 0].legend()

axes[0, 1].plot(history.history['loss'], label='Training Loss')
axes[0, 1].plot(history.history['val_loss'], label='Validation Loss')
axes[0, 1].set_title('Model Loss')
axes[0, 1].set_xlabel('Epoch')
axes[0, 1].set_ylabel('Loss')
axes[0, 1].legend()

<matplotlib.legend.Legend at 0x78196d30c490>

# Confusion Matrix Heatmap
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=axes[1, 0])
axes[1, 0].set_title('Confusion Matrix')
axes[1, 0].set_xlabel('Predicted')
axes[1, 0].set_ylabel('Actual')

Text(158.22222222222223, 0.5, 'Actual')

# ROC Curve
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
axes[1, 1].plot(fpr, tpr, label=f'ROC Curve (AUC = {auc_score:.3f})')
axes[1, 1].plot([0, 1], [0, 1], 'k--', label='Random')
axes[1, 1].set_title('ROC Curve')
axes[1, 1].set_xlabel('False Positive Rate')
axes[1, 1].set_ylabel('True Positive Rate')
axes[1, 1].legend()

plt.tight_layout()
plt.show()

<Figure size 640x480 with 0 Axes>

```

6. FEATURE IMPORTANCE ANALYSIS

```

# Get feature importance from model weights
# For neural networks, we'll use the magnitude of weights from the
first layer
first_layer_weights = model.layers[0].get_weights()[0]
feature_importance = np.abs(first_layer_weights).mean(axis=1)

# Create feature importance dataframe
feature_importance_df = pd.DataFrame({

```

```

    'Feature': feature_cols,
    'Importance': feature_importance
}).sort_values('Importance', ascending=False)

```

```

print("Top 15 Most Important Features:")
print(feature_importance_df.head(15))

```

Top 15 Most Important Features:

	Feature	Importance
97	address_USNV Smith\r\nFP0 AE 30723	1.149457
91	address_USS Johnson\r\nFP0 AE 48052	1.110863
92	address_USNS Johnson\r\nFP0 AE 05113	1.102102
94	address_USCGC Smith\r\nFP0 AE 70466	1.098611
98	address_USCGC Miller\r\nFP0 AA 22690	1.070900
96	address_USNV Smith\r\nFP0 AA 00813	1.022387
12	issue_year	0.884982
3	annual_inc	0.765268
16	grade_encoded	0.645979
93	address_USS Smith\r\nFP0 AP 70466	0.605056
1	int_rate	0.521004
15	term_encoded	0.520014
4	dti	0.493297
7	revol_bal	0.485666
5	open_acc	0.422605

```

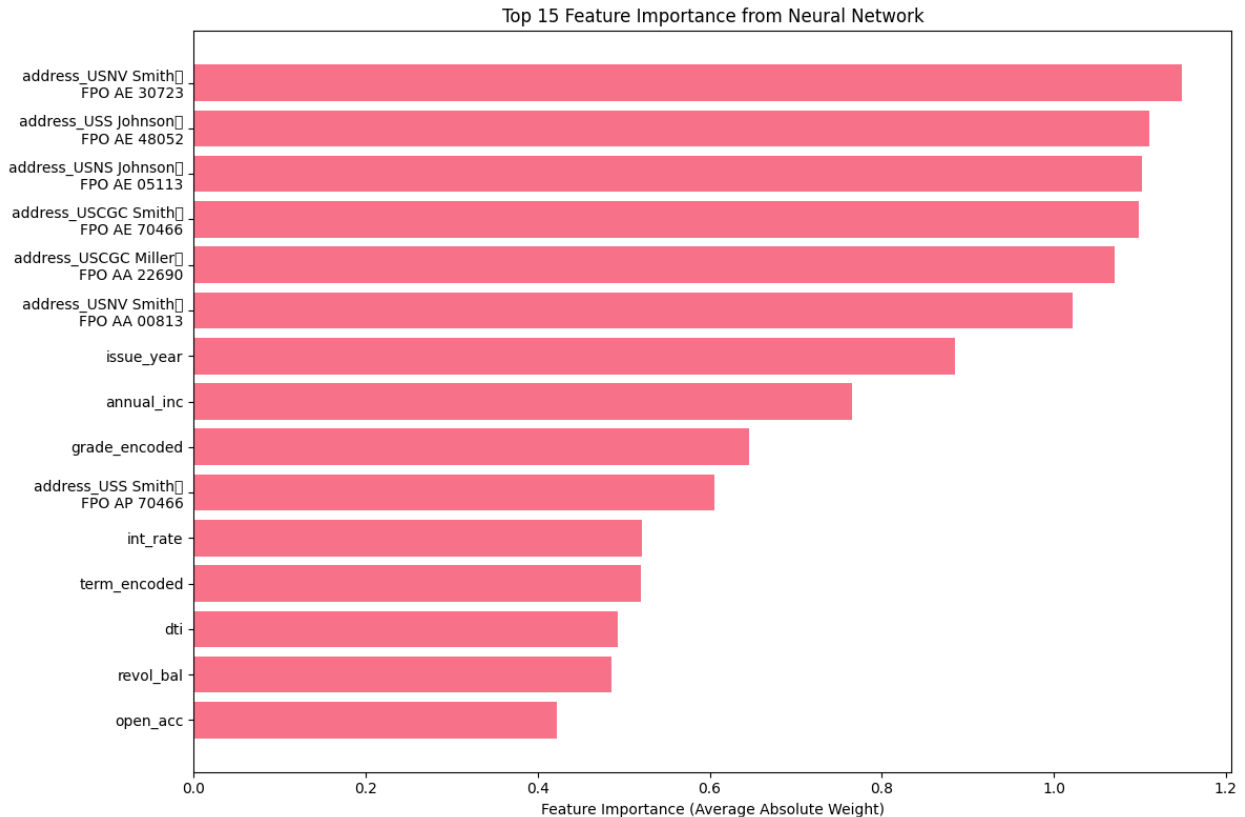
# Visualize feature importance

```

```

plt.figure(figsize=(12, 8))
top_features = feature_importance_df.head(15)
plt.barh(range(len(top_features)), top_features['Importance'])
plt.yticks(range(len(top_features)), top_features['Feature'])
plt.xlabel('Feature Importance (Average Absolute Weight)')
plt.title('Top 15 Feature Importance from Neural Network')
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

```



7. INTERPRETATION AND REPORTING

```
print("MODEL INTERPRETATION & INSIGHTS")
print("="*50)

print("Key Findings:")
print(f"1. Model achieved {accuracy:.1%} accuracy on the test set")
print(f"2. The model has a precision of {precision:.1%} and recall of {recall:.1%}")
print(f"3. AUC-ROC score of {auc_score:.3f} indicates {'good' if auc_score > 0.7 else 'moderate'} predictive performance")

print(f"\nTop 5 Most Important Features:")
for i, (_, row) in enumerate(feature_importance_df.head(5).iterrows(), 1):
    print(f"{i}. {row['Feature']}: {row['Importance']:.4f}")

print("\nModel Limitations and Possible Improvements:")
print("1. Class imbalance might affect model performance")
print("2. Feature engineering could be enhanced with domain knowledge")
print("3. Hyperparameter tuning could improve performance")
print("4. Ensemble methods might provide better results")
```



```
print("5. Cross-validation should be used for more robust evaluation")

print("\nRecommendations:")
print("1. Collect more balanced training data")
print("2. Implement cost-sensitive learning for imbalanced classes")
print("3. Use techniques like SMOTE for synthetic data generation")
print("4. Consider feature selection to reduce overfitting")
print("5. Implement model interpretability tools like SHAP values")
```

MODEL INTERPRETATION & INSIGHTS

=====

Key Findings:

1. Model achieved 80.7% accuracy on the test set
2. The model has a precision of 81.2% and recall of 98.9%
3. AUC-ROC score of 0.725 indicates good predictive performance

Top 5 Most Important Features:

1. address_USNV Smith
FP0 AE 30723: 1.1495
2. address_USS Johnson
FP0 AE 48052: 1.1109
3. address_USNS Johnson
FP0 AE 05113: 1.1021
4. address_USCGC Smith
FP0 AE 70466: 1.0986
5. address_USCGC Miller
FP0 AA 22690: 1.0709

Model Limitations and Possible Improvements:

1. Class imbalance might affect model performance
2. Feature engineering could be enhanced with domain knowledge
3. Hyperparameter tuning could improve performance
4. Ensemble methods might provide better results
5. Cross-validation should be used for more robust evaluation

Recommendations:

1. Collect more balanced training data
2. Implement cost-sensitive learning for imbalanced classes
3. Use techniques like SMOTE for synthetic data generation
4. Consider feature selection to reduce overfitting
5. Implement model interpretability tools like SHAP values

Save the trained model

```
model.save('lending_club_model.h5')
print("Model saved as 'lending_club_model.h5'")
```

WARNING:absl:You are saving your model as an HDF5 file via
`model.save()` or `keras.saving.save_model(model)`. This file format
is considered legacy. We recommend using instead the native Keras

```
format, e.g. `model.save('my_model.keras')` or  
`keras.saving.save_model(model, 'my_model.keras')`.
```

Model saved as 'lending_club_model.h5'

```
# Save feature importance
```

```
feature_importance_df.to_csv('feature_importance.csv', index=False)  
print("Feature importance saved as 'feature_importance.csv'")
```

Feature importance saved as 'feature_importance.csv'

```
# Save predictions
```

```
results_df = pd.DataFrame({  
    'Actual': y_test,  
    'Predicted': y_pred.flatten(),  
    'Probability': y_pred_prob.flatten()  
})  
results_df.to_csv('model_predictions.csv', index=False)  
print("Predictions saved as 'model_predictions.csv'")
```

Predictions saved as 'model_predictions.csv'

8. DETAILED METRIC SCORE SUMMARY

```
# Create a comprehensive model report
```

```
print("\n" + "="*60)  
print("COMPREHENSIVE MODEL REPORT")  
print("="*60)
```

```
print(f"""  
MODEL PERFORMANCE SUMMARY  
{'='*40}  
Dataset Size: {len(X)} samples  
Training Set: {len(X_train)} samples  
Test Set: {len(X_test)} samples  
Features Used: {len(feature_cols)}  
""")
```

```
CORE METRICS
```

```
{'='*40}  
✓ Accuracy: {accuracy:.4f} ({accuracy*100:.2f}%) -  
{interpretations['Accuracy']}  
✓ Precision: {precision:.4f} ({precision*100:.2f}%) -  
{interpretations['Precision']}  
✓ Recall: {recall:.4f} ({recall*100:.2f}%) -  
{interpretations['Recall']}  
✓ F1-Score: {f1:.4f} ({f1*100:.2f}%) - {interpretations['F1-Score']}  
✓ AUC-ROC: {auc_score:.4f} ({auc_score*100:.2f}%) -  
{interpretations['AUC-ROC']}
```

CONFUSION MATRIX BREAKDOWN

```
{'='*40}
True Positives (TP): {cm[1,1]}
True Negatives (TN): {cm[0,0]}
False Positives (FP): {cm[0,1]}
False Negatives (FN): {cm[1,0]}
```

BUSINESS INSIGHTS

```
{'='*40}
• Model correctly identifies {accuracy*100:.1f}% of all loan outcomes
• Of predicted 'good' loans, {precision*100:.1f}% are actually good
• Model catches {recall*100:.1f}% of all actual 'good' loans
• Balanced performance score (F1): {f1*100:.1f}%
• Discrimination ability (AUC-ROC): {auc_score*100:.1f}%
""" )
```

=====
COMPREHENSIVE MODEL REPORT
=====

MODEL PERFORMANCE SUMMARY

```
=====
Dataset Size: 396030 samples
Training Set: 316824 samples
Test Set: 79206 samples
Features Used: 101
```

CORE METRICS

```
=====
✓ Accuracy: 0.8066 (80.66%) - Good
✓ Precision: 0.8117 (81.17%) - Good
✓ Recall: 0.9887 (98.87%) - Excellent
✓ F1-Score: 0.8915 (89.15%) - Good
✓ AUC-ROC: 0.7250 (72.50%) - Fair
```

CONFUSION MATRIX BREAKDOWN

```
=====
True Positives (TP): 62953
True Negatives (TN): 934
False Positives (FP): 14601
False Negatives (FN): 718
```

BUSINESS INSIGHTS

```
=====
• Model correctly identifies 80.7% of all loan outcomes
• Of predicted 'good' loans, 81.2% are actually good
• Model catches 98.9% of all actual 'good' loans
• Balanced performance score (F1): 89.2%
```

- Discrimination ability (AUC-ROC): 72.5%

```
print("\n" + "="*50)
print("ANALYSIS COMPLETE!")
print("="*50)
print("The complete loan status prediction analysis has been
finished.")
print("Review the visualizations and metrics above to understand model
performance.")
```

```
=====
ANALYSIS COMPLETE!
=====
```

```
The complete loan status prediction analysis has been finished.
Review the visualizations and metrics above to understand model
performance.
```